



UNIVERSITY of LIMERICK

O L L S C O I L L U I M N I G H

FACULTY of SCIENCE and ENGINEERING

Department of Computer Science
and Information Systems

Mid-Term Assessment Paper

Academic Year: 2021-2022

Module Title: Deep Reinforcement Learning

Duration of Exam: 1 Hour

Lecturer(s): J.J. Collins

Semester: Spring

Module Code: CS6482

Percent of Total Marks: 10

Paper marked out of : 10

11:00-12:00 Thur. 24th February 2022 (Week 5)

Instructions to Candidates:

- Answer ALL ten questions.
- All questions carry equal marks.
- NO RED PEN

Name: _____

ID Number: _____

Course: _____

Q1. Describe the Symbolic AI and Machine Learning paradigms, and illustrate the answer with examples.

Symbolic AI is

Machine Learning is

Q2: What is the formula for the update applied to a weight in a hidden layer when performing Back Propagation (BP) in a Multi-Layer Perceptron (MLP)?

Q3: The following code fragments are from Nielsen's implementation of an MLP for MNIST.

```
L1     ##### network.py
L2     class Network(object):

L3         def __init__(self, sizes):
L4             self.num_layers = len(sizes)
L5             self.sizes = sizes
L6             self.biases = [np.random.randn(y, 1) for y in sizes[1:]]
L7             self.weights = [np.random.randn(y, x) for x, y in zip(sizes[:-1], sizes[1:])]

L8     ##### test.py
L9     import network
L10    net = network.Network([784, 30, 10])
```

What are the dimensions of the three lists

1. sizes,
2. sizes[1:]
3. zip(sizes[:-1], sizes[1:])

Q4: What is a vanishing gradient? What causes vanishing gradient? And what technique(s) can be used to reduce the impact of vanishing gradients? Include diagram(s) if discussing activation functions.

Q5: Explain why the number of parameters in GoogleLeNet using Inception modules is significantly less than AlexNet - 6 million v of 60 million. The answer should focus exclusively on the Inception module. Illustrate the answer with a diagram and/or calculations.

Q6: Describe the key concept(s) in ResNet. Include a discussion on the purpose of a kernel of size 1×1 with stride 2. Illustrate the discussion with a diagram.

Q7: Given an input image of dimensions 75×100 with 3 channels for RGB. And a convolutional layer with 5×5 kernels stride 1, outputting 200 feature maps of size 75×100 . How many parameters in the convolutional layer? Please show your calculations.

Q8: Briefly describe four techniques that can be used to reduce overfitting

1

2

3

4

Q9: Describe the two key concerns addressed in Reinforcement Learning paradigms?

1

2

Q10: What happens in Figure 1 when the number of steps approaches infinity?




Figure 1: Sutton and Barto. Introduction to Reinforcement Learning, 2nd Edition. The MIT Press. 2018